Task-Relative Descriptive Privilege in Linear Gaussian Dynamical Systems

Kunal Bhatia [1]

[1] *Independent Researcher, Heidelberg, Germany*

(Dated: April 19, 2026)

In a finite-dimensional linear Gaussian dynamical system, the rank-r linear observer minimizing squared-error prediction loss is uniquely determined—up to rotation within the eigenspace—by the top-r eigenspace of the target information matrix. Two prediction targets with distinct dominant eigenspaces admit no jointly optimal observer. Global privilege—a single observer simultaneously optimal for all tasks—holds exactly when the target family is *r-coherent*: every information matrix shares a common top-r eigenspace. This is an exact algebraic condition. A two-target instance with $n = 8$, $r = 2$ yields a privilege gap $\Gamma = 0.026 > 0$; a structurally coherent construction attains zero pairwise subspace angle and is perturbation-robust.

I. INTRODUCTION

A basic question in the study of dynamical systems is whether there is a single low-dimensional description that is optimal for every prediction task, or whether the optimal description depends on the target one seeks to predict. This paper answers that question exactly for finite-dimensional linear Gaussian dynamical systems (LGDS). The optimal rank-r observer is the top-r eigenspace of a target information matrix. Two targets whose information matrices have distinct dominant subspaces cannot be simultaneously optimised by any single observer. The exception has an exact algebraic characterisation: r-coherence of the target family.

The linear state-space framework is classical [1]; the relevant matrix optimisation on the Grassmannian [2] is governed by the Ky Fan maximum principle [3].

This yields three things useful for the present question: an exact risk formula with no approximation, a proof that task-relative privilege is the rule within this class, and an explicit algebraic criterion for when global privilege is warranted.

This result is intentionally narrow. It concerns only finite-dimensional linear Gaussian dynamical systems, fixed-rank linear observers, and squared-error prediction targets.

II. SYSTEM AND OBSERVER FAMILY

a. Dynamics. Let

$$x_{t+1} = Ax_t + \xi_t, \quad \xi_t \sim \mathcal{N}(0, Q), \quad x_t \in \mathbb{R}^n, \tag{1}$$

with A stable (spectral radius < 1) and $Q \succ 0$. The unique stationary covariance Σ satisfies $A\Sigma A^\top + Q = \Sigma$ and is strictly positive definite [1].

b. Observer family. An observer of rank r ($1 \le r < n$) is parametrised by a row-orthonormal matrix $U \in \mathbb{R}^{r\times n}$, $UU^\top = I_r$:

$$\phi_U(x) = U\Sigma^{-1/2}x. \tag{2}$$

Pre-whitening by $\Sigma^{-1/2}$ places all observers on a common isotropic footing; in the static case ($A = I$, $\tau = 0$) this reduces to principal component analysis [4]. Two observers differing only by an orthogonal rotation of their row spaces are equivalent; the observer space is the Grassmannian $\mathrm{Gr}(r, n)$, a smooth compact Riemannian manifold of dimension $r(n-r)$, with metric and geodesic structure from $O(n)$ [2].

c. Target family. A prediction target is a pair (C, τ) with $C \in \mathbb{R}^{p\times n}$ and integer $\tau \ge 1$:

$$y_{C,\tau} = Cx_{t+\tau}. \tag{3}$$

Performance is total mean squared prediction error. A target family $\mathcal{T} = \{(C_k, \tau_k)\}$ is any finite collection.

III. THEOREM

A. Exact Bayes risk formula

Given observation $z = U\Sigma^{-1/2}x_t$, the MMSE predictor of $Cx_{t+\tau}$ is linear by Gaussian conjugacy. At stationarity, $\mathrm{Cov}(Cx_{t+\tau}, z^\top) = CA^\tau\Sigma^{1/2}U^\top$ and $\mathrm{Cov}(z, z^\top) = I_r$. The Bayes risk is

$$\mathcal{R}(U;\, C, \tau) = \mathrm{tr}(C\Sigma C^\top) - \mathrm{tr}\big(U\, M_{C,\tau}\, U^\top\big), \tag{4}$$

where the *target information matrix* is

$$M_{C,\tau} = \Sigma^{1/2}(A^\tau)^\top C^\top CA^\tau\Sigma^{1/2} \in \mathbb{R}^{n\times n}. \tag{5}$$

$M_{C,\tau}$ is symmetric positive semidefinite of rank at most $\min(p, n)$. The formula is exact.

B. Top-eigenspace optimality

Proposition 1. *For fixed rank r, the observer minimising $\mathcal{R}(U;\, C, \tau)$ over $\mathrm{Gr}(r, n)$ has row space equal to the top-r eigenspace of $M_{C,\tau}$. The minimum risk is*

$$\mathcal{R}^* = \mathrm{tr}(C\Sigma C^\top) - \sum_{i=1}^{r} \lambda_i(M_{C,\tau}), \tag{6}$$

where $\lambda_1 \ge \cdots \ge \lambda_n$ are the eigenvalues of $M_{C,\tau}$. Uniqueness up to rotation holds when $\lambda_r > \lambda_{r+1}$.

Proof. Minimising $\mathcal{R}$ is equivalent to maximising $\mathrm{tr}(UM_{C,\tau}U^\top)$ subject to $UU^\top = I_r$. By the Ky Fan variational principle [3, 5] the maximum equals $\sum_{i=1}^{r} \lambda_i(M_{C,\tau})$, achieved exactly when the row space of U spans the corresponding eigenspace. ■

C. No-global-privilege corollary

Corollary 2. *If targets (C_1, τ_1) and (C_2, τ_2) have information matrices M_1 and M_2 with distinct top-r eigenspaces, no single rank-r observer is simultaneously optimal for both.*

Proof. Optimality for target k requires the row space of U to equal the top-r eigenspace of M_k. Distinct eigenspaces cannot be spanned by the same U. ■

The corollary extends to any finite target family: if $\{M_k\}$ do not share a common top-r eigenspace, no rank-r observer is globally optimal.

D. r-coherence: the exact exception

Definition 3. *A target family $\mathcal{T}$ is* r-coherent *with respect to (A, Σ) if all matrices $\{M_{C,\tau} : (C, \tau) \in \mathcal{T}\}$ share a common top-r eigenspace V^*.*

Corollary 4. *Global privilege holds if and only if $\mathcal{T}$ is r-coherent.*

r-coherence requires all targets to be linear functionals of a common r-dimensional sufficient statistic across all horizons—a strong structural alignment between the dynamics and the task family. Outside r-coherence, the optimal observer generally depends on the target within the class studied here.

Remark 5. *r-coherence is an algebraic equality between eigenspaces, not an approximate or asymptotic condition. It is the exact characterisation of when global privilege is warranted.*

IV. NUMERICAL ILLUSTRATIONS

The theorem is analytically exact under its stated assumptions. We illustrate it numerically with two finite-dimensional instances: a no-global-privilege example and a structurally coherent example. All reported quantities are computed exactly from the analytic Bayes risk formula and numerical optimization on the Grassmannian.

a. Setup. State dimension $n = 8$, observer rank $r = 2$. The model is a finite-dimensional linear Gaussian dynamical system [6]. $A = VDV^{-1}$ with eigenvalues drawn from $U[0.3, 0.85]$, random non-orthogonal V (seed 100), $Q = 0.1\, I_8$. Σ is computed from the discrete Lyapunov equation; $\Sigma \succ 0$ is confirmed to machine precision. All reported quantities are computed exactly from the analytic Bayes risk formula, eigenspace calculations, and repeated projected optimization over $\mathrm{Gr}(2, 8)$.

TABLE I. Regret matrix, Condition A ($n = 8$, $r = 2$, seed 100). All values from the exact analytic Bayes risk formula.

Observer	Regret on T_1	Regret on T_2
U_1^* (optimal for T_1)	0.000	0.036
U_2^* (optimal for T_2)	0.763	0.000
U_{joint} (compromise)	0.026	0.026

b. Condition A: no-global-privilege instance. Targets $T_1 = ([e_1^\top; e_2^\top], 1)$ and $T_2 = ([e_5^\top; e_6^\top], 5)$.

With non-normal A, A^1 and A^5 rotate state-space directions differently; the principal angles [7] between the top-2 eigenspaces of M_1 and M_2 are 54.1° and 24.0°, confirming that no single rank-2 observer is simultaneously optimal.

The global-privilege gap

$$\Gamma = \min_{U \in \mathrm{Gr}(r,n)} \max_k \left[\mathcal{R}(U; C_k, \tau_k) - \mathcal{R}_k^*\right] \tag{7}$$

measures the best achievable worst-case regret under task conflict. The minimax observer U_{joint}, found by repeated projected optimization over $\mathrm{Gr}(2, 8)$ from random initializations, achieves $\Gamma = 0.026 > 0$. No globally optimal observer exists for this target family. Table I gives the full regret structure. The asymmetry in the off-diagonal entries reflects the geometry of the instance: U_2^* is optimised for directions nearly orthogonal to T_1, so its cross-regret is large; U_1^* retains partial overlap with the long-horizon T_2 target, so its cross-regret is small. Both are computed from the same analytic formula.

c. Condition B′: structural coherence and its robustness. A target family is constructed so that all M_k share the same top-2 eigenspace by design: normal $A_{\text{sym}} = VDV^\top$ with orthogonal V, and a common C whose rows are the top-2 eigenvectors of A, varying only the prediction horizon $\tau \in \{1, 2, 3, 5\}$. Under this construction $M_{C,\tau}$ has the same dominant 2-dimensional eigenspace for all τ, so the four optimal observers are identical by the theorem. The maximum pairwise principal angle across all four is 0.000° to machine precision.

Under $\varepsilon = 0.1$ perturbations to C (20 random draws), the maximum principal angle across all draws and all pairs is 0.102 rad—well below the fragility threshold of 0.15 rad. Structural coherence is perturbation-robust, distinguishing it from accidental near-coincidences of eigenspaces that break under small perturbations.

The numerical results are consistent with the theorem in all cases.

V. SCOPE

a. What is earned. In finite-dimensional linear Gaussian dynamical systems at stationarity, among rank-r

linear observers $\phi_U(x) = U\Sigma^{-1/2}x$ and squared-error prediction targets $y = Cx_{t+\tau}$, descriptive privilege is task-relative unless the target family is r-coherent. r-coherence requires strong structural alignment between dynamics and tasks. Outside it, no fixed-rank observer is globally optimal, and the cost of forcing any observer onto a non-native target is given exactly by the Bayes risk formula. The Grassmannian $\mathrm{Gr}(r, n)$ is the literal domain of this optimisation: a compact Riemannian manifold on which observer quality is a smooth function, with $O(n)$ acting by transformation.

b. What is not earned. This result does not apply to nonlinear systems, non-Gaussian noise, nonlinear observers, or non-prediction tasks. The causal axis—whether the observer optimal for intervention differs from the observer optimal for prediction—is not addressed here.

ACKNOWLEDGMENTS

No external funding was received for this work. Code to reproduce all numerical illustrations is available at https://github.com/kunalb541/LGDS.

[1] T. Kailath, *Linear Systems* (Prentice-Hall, Englewood Cliffs, NJ, 1980).
[2] A. Edelman, T. A. Arias, and S. T. Smith, SIAM Journal on Matrix Analysis and Applications **20**, 303 (1998).
[3] K. Fan, Proceedings of the National Academy of Sciences of the United States of America **35**, 652 (1949).
[4] H. Hotelling, Journal of Educational Psychology **24**, 417 (1933).
[5] R. A. Horn and C. R. Johnson, *Matrix Analysis*, 2nd ed. (Cambridge University Press, Cambridge, 2013).
[6] R. E. Kalman, Journal of Basic Engineering **82**, 35 (1960).
[7] Å. Björck and G. H. Golub, Mathematics of Computation **27**, 579 (1973).